
Not All Tokens Are Equally Useful for Steering: Robust Directions and Prefix Steering

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Activation steering controls language model behavior by modifying hidden-state ac-
2 tivations along interpretable concept directions, often constructed using Difference-
3 in-Means (DiM) over contrastive examples. However, recent work has observed
4 that DiM steering can degrade general capabilities. In this work, we revisit activa-
5 tion steering from two complementary perspectives: steering-direction construction
6 and intervention-time token selection. First, we show that DiM should not be
7 viewed purely as a variance-reduction problem achieved by averaging more activa-
8 tions. Both theoretical analysis and empirical results demonstrate that increasing
9 the number of activations, either by using more prompt and response tokens or
10 by increasing the number of contrastive examples, does not necessarily improve
11 steering performance and can even degrade it. We identify the key challenge as
12 interference from stable but intervention-irrelevant concepts encoded in hidden
13 states, motivating a reinterpretation of DiM as a robust contrastive mean estimation
14 problem. Based on this perspective, we propose **Robust DiM**, which estimates
15 contrastive directions using activations that are both representative of their own
16 class and discriminative from the opposite class. Across three target concepts and
17 four models, Robust DiM improves target-concept steering by up to 8% while
18 improving general-capability preservation by up to 5% over standard DiM. Sec-
19 ond, we revisit the common practice of applying steering interventions at every
20 generation step. We show that most steering effects emerge within the first few
21 generated tokens, while later interventions primarily accumulate distributional shift
22 and degrade general capability. Motivated by this observation, we propose **prefix**
23 **steering**, which applies steering updates only to an initial generation prefix. Prefix
24 steering recovers most of the target-control effect of full-generation intervention
25 while substantially reducing interference with general capabilities. Together, our
26 results suggest that effective activation steering depends critically on selecting both
27 intervention-relevant construction activations and intervention-relevant generation
28 positions, rather than uniformly averaging or steering across all tokens.

1 Introduction

30 Large language models (LLMs) exhibit strong capabilities across a wide range of tasks [1–3], but
31 understanding and controlling their internal representations remains a central challenge [4–6]. A
32 growing line of work addresses this problem by intervening directly on hidden activations, using
33 representation-level directions to elicit, suppress, or modify model behaviors [7–9]. Among these
34 methods, Difference-in-Means (DiM) is especially attractive: given positive and negative examples
35 of a target behavior, it subtracts the mean hidden state of the negative set from that of the positive
36 set, and uses the resulting direction for activation steering [10–12]. Despite its simplicity, DiM often
37 outperforms other more complicated approaches such as sparse autoencoders [13]. However, recent

work has also shown that activation steering can interfere with unrelated model capabilities, causing degradation in fluency, reasoning, or benchmark performance [14–17]. Although these trade-offs are increasingly recognized empirically, the underlying causes remain poorly understood.

In this work, we revisit activation steering from two complementary perspectives:

- Q1: Direction construction.** Which activations should be used to construct the steering direction?
- Q2: Intervention placement.** Which tokens should receive steering interventions during generation?

The standard intuition of DiM is that averaging many activations cancels variation unrelated to the target behavior while preserving the concept-relevant signal [18–20]. Under this view, using more activations should produce a cleaner, and thus more effective steering direction. Our starting point is the observation that hidden-state activations encode many simultaneous concepts beyond the target behavior, including topic, syntax, stylistic patterns, instruction formats, and other context-dependent information. Consequently, averaging over more activations does not necessarily improve the quality of the identified concept direction. Instead, additional activations can introduce stable but intervention-irrelevant structure that contaminates the estimated direction, creating a gap between directional stability and steering effectiveness. We observe this gap consistently across various models and target concepts: averaging all prompt-token activations can underperform averaging only the last few prompt tokens, prompt-only construction can outperform prompt-response construction, and increasing the number of contrastive examples can improve directional stability without improving downstream steering. These results suggest that many token activations contain statistically stable but behaviorally irrelevant information.

Motivated by this diagnosis, we introduce **Robust DiM**, a token-selective direction-construction method. Rather than treating all candidate activations as exchangeable, Robust DiM estimates contrastive directions from activations that are both representative of their own class and discriminative from the opposite class. Representative activations reduce within-class nuisance variation, while discriminative activations preserve the separation needed for behavioral control. Across three target concepts and four models, Robust DiM improves target-concept steering by up to 8% and improves general-capability preservation by up to 5% over standard DiM.

We then revisit a second, largely overlooked assumption in activation steering: the standard practice of applying steering interventions at every generation step [7, 11, 12]. While this uniform intervention policy is simple, it implicitly assumes that all generated token positions are equally useful intervention sites. Our experiments through KL divergence and steering performance show instead that most steering effects emerge within the first few generated tokens, whereas later interventions primarily accumulate distributional shift and interfere with unrelated computation. Based on this observation, we introduce **prefix steering**, which applies steering updates only to an initial prefix of generated tokens. Prefix steering retains most of the steering gain achieved by all-token intervention, and in some cases matches it, while substantially better preserving general capability (improving MMLU-Pro performance by up to 6%).

Our main contributions can be summarized as follows:

- We revisit the statistical assumptions underlying DiM and use the resulting analysis to shed light on several practical design choices, including token position, activation source, and the number of contrastive examples. Our results show that concept direction is better viewed as a robust contrastive mean estimation problem rather than a pure variance-reduction problem through averaging.
- We introduce **Robust DiM**, which selects discriminative activations for contrastive direction construction. Across three target concepts and four models, Robust DiM improves target-concept steering by up to 8% and improves general-capability preservation by up to 5% over standard DiM.
- We introduce **prefix steering**, which applies the steering update only to early generated tokens. Prefix steering recovers most of the steering gain of full-generation intervention while substantially reducing interference with general capabilities (improving MMLU-Pro performance by up to 6%).

2 Activation Steering Preliminaries

Transformer Decoder-only transformers [21] map input tokens $\mathbf{t} = (t_1, t_2, \dots, t_n) \in \mathcal{V}^n$ to output probability distributions over the vocabulary $\mathcal{V}$. Let $\mathbf{h}_i^{(\ell)} \in \mathbb{R}^d$ denote the residual stream activation of the i -th token at the ℓ -th layer, where d is the dimensionality of the hidden state. Each of the L

transformer layers applies a sequence of attention and MLP transformations to update the residual stream:

$$\tilde{\mathbf{h}}_i^{(\ell)} = \mathbf{h}_i^{(\ell)} + \text{Attn}^{(\ell)}(\mathbf{h}_{1:i}^{(\ell)}), \quad \mathbf{h}_i^{(\ell+1)} = \tilde{\mathbf{h}}_i^{(\ell)} + \text{MLP}^{(\ell)}(\tilde{\mathbf{h}}_i^{(\ell)}). \quad (1)$$

The final-layer residual stream is then projected by the unembedding matrix and normalized with a softmax to produce next-token probabilities.

Linear representation hypothesis. A central assumption underlying many representation-steering methods is the *linear representation hypothesis* [18], which posits that high-level, semantically meaningful concepts are approximately encoded as linear directions in activation space¹. For notational simplicity, we omit the layer index when considering activations at a given layer, writing $\mathbf{h}_i$ for the residual-stream activation of the i -th token. Under this view, $\mathbf{h}_i$ can be written as

$$\mathbf{h}_i = \sum_j \alpha_{i,j} \mathbf{r}_j + \boldsymbol{\epsilon}_i, \quad (2)$$

where $\mathbf{r}_j$ denotes a concept or feature direction, $\alpha_{i,j}$ is the corresponding activation coefficient or strength at token position i , and $\boldsymbol{\epsilon}_i$ captures residual variation not explained by this linear decomposition. This perspective motivates methods for understanding, monitoring, and controlling model behavior through interventions on internal representations.

Contrastive direction construction. Given the linear-direction framework, a key challenge lies in estimating the specific direction that corresponds to a target behavior. Various methods have been proposed to isolate such concept directions [24–26]. Among them, Difference-in-Means (DiM) is a simple and widely used approach that compares activations induced by positive and negative examples [27–29]. Recent studies have shown that DiM outperforms more complex approaches such as sparse autoencoders [13]. Let $\mathcal{D}^+$ and $\mathcal{D}^-$ denote datasets corresponding to two contrastive classes for a target concept or behavior. For a given layer, we collect token activations from the two datasets and denote them by

$$\mathcal{A}^+ = \{\mathbf{h}_i(x) : x \in \mathcal{D}^+, i \in \mathcal{S}(x)\}, \quad \mathcal{A}^- = \{\mathbf{h}_i(x) : x \in \mathcal{D}^-, i \in \mathcal{S}(x)\}, \quad (3)$$

where $\mathcal{S}(x)$ is a token-selection rule specifying which token positions of example x are used to construct the direction. The positive and negative mean activations are then computed as

$$\boldsymbol{\mu}^+ = \frac{1}{|\mathcal{A}^+|} \sum_{\mathbf{h} \in \mathcal{A}^+} \mathbf{h}, \quad \boldsymbol{\mu}^- = \frac{1}{|\mathcal{A}^-|} \sum_{\mathbf{h} \in \mathcal{A}^-} \mathbf{h}. \quad (4)$$

The DiM direction is $\mathbf{r} = \boldsymbol{\mu}^+ - \boldsymbol{\mu}^-$. Following prior work [30], we use the normalized direction $\hat{\mathbf{r}} = \frac{\mathbf{r}}{\|\mathbf{r}\|_2}$, so that the steering strength is controlled only by a scalar intervention coefficient rather than by the norm of the estimated direction.

Activation steering. Given an identified concept direction $\hat{\mathbf{r}}$, activation steering modifies the hidden-state activations during autoregressive generation. Let $\mathbf{h}_k$ denote the hidden-state activation corresponding to the k -th generated token at a given layer. At each generation step, we can either elicit the target behavior by adding the steering direction, or suppress the target behavior by ablating the activation component along that direction:

$$\mathbf{h}_k \leftarrow \mathbf{h}_k + a\hat{\mathbf{r}} \quad \text{or} \quad \mathbf{h}_k \leftarrow \mathbf{h}_k - \beta\hat{\mathbf{r}}\hat{\mathbf{r}}^\top \mathbf{h}_k. \quad (5)$$

Here $a \geq 0$ controls the additive steering strength, while b controls the ablation strength. Since $\hat{\mathbf{r}}$ points from the negative-class mean toward the positive-class mean, the additive update shifts generated-token activations toward the target concept. When $\beta = 1$, the ablation update fully removes the component of $\mathbf{h}_k$ along $\hat{\mathbf{r}}$. In the standard formulation, this intervention is applied at every generation step, i.e., to the hidden-state activation of every generated token during inference. This implicitly treats all generated token positions as equally useful intervention sites.

3 Robust Difference-in-Means

While DiM has been widely used in practice, its statistical foundations remain relatively underexplored, with the notable exception of the pioneering work of [31], which studies a simplified single-concept Gaussian setting. In this section, we revisit the statistical assumptions underlying DiM and use the resulting analysis to shed light on several practical design choices, including the

¹While some concepts may exhibit non-linear structures or lack a singular global direction [19, 22, 23], linear directions frequently serve as effective approximations for behaviorally relevant features.

prompt-token positions and activation sources used for constructing steering directions, as well as the role of the number of contrastive examples. Our analysis and experiments suggest that DiM should be viewed not merely as a variance-reduction problem through averaging, but rather as a robust contrastive mean estimation problem under heterogeneous and interfering concepts. Motivated by this perspective, we then develop a Robust DiM approach that retains the simplicity of standard DiM while more reliably identifying intervention-relevant concept directions.

3.1 Rethinking the Statistical Assumptions Behind DiM

To begin, let $\mathbf{r}_c$ denote the direction associated with the target concept. We can decompose (2) as

$$\mathbf{h}_i = \alpha_{i,c} \mathbf{r}_c + \sum_{j \neq c} \alpha_{i,j} \mathbf{r}_j + \epsilon_i. \quad (6)$$

Let $\bar{\alpha}_j^\pm$ and $\bar{\epsilon}^\pm$ denote the averages of the concept coefficients $\{\alpha_{i,j}\}_i$ and residuals ϵ_i over $\mathcal{A}^\pm$, respectively. We further define the corresponding differences in means as $\Delta\alpha_j$ and $\Delta\epsilon$:

$$\bar{\alpha}_j^\pm = \frac{1}{|\mathcal{A}^\pm|} \sum_i \alpha_{i,j}^\pm, \quad \bar{\epsilon}^\pm = \frac{1}{|\mathcal{A}^\pm|} \sum_i \epsilon_i^\pm, \quad \Delta\alpha_j = \bar{\alpha}_j^+ - \bar{\alpha}_j^-, \quad \Delta\epsilon = \bar{\epsilon}^+ - \bar{\epsilon}^-. \quad (7)$$

Substituting this decomposition into the DiM direction $\mathbf{r} = \mu^+ - \mu^-$ gives

$$\mathbf{r} = \frac{1}{|\mathcal{A}^+|} \sum_{\mathbf{h}_i \in \mathcal{A}^+} \mathbf{h}_i - \frac{1}{|\mathcal{A}^-|} \sum_{\mathbf{h}_i \in \mathcal{A}^-} \mathbf{h}_i = \underbrace{\frac{\Delta\alpha_c \mathbf{r}_c}{\text{target concept}}}_{\text{target concept}} + \underbrace{\sum_{j \neq c} \Delta\alpha_j \mathbf{r}_j}_{\text{non-target concepts}} + \underbrace{\Delta\epsilon}_{\text{residual}}. \quad (8)$$

We now characterize the alignment between the DiM direction $\mathbf{r}$ and the target concept direction $\mathbf{r}_c$.

Lemma 1. *Suppose that the concept directions $\mathbf{r}_j$ are orthogonal and have unit norm, and the residual terms ϵ_i are i.i.d. Gaussian random vectors with zero mean. Then, with high probability, the correlation between the DiM direction $\mathbf{r}$ and the concept direction $\mathbf{r}_c$ satisfies*

$$\cos(\mathbf{r}, \mathbf{r}_c) = \frac{\Delta\alpha_c}{\sqrt{(\Delta\alpha_c)^2 + \sum_{j \neq c} (\Delta\alpha_j)^2}} + O\left(\frac{1}{\sqrt{|\mathcal{A}^+|}}\right). \quad (9)$$

Proof. Since each ϵ_i is an iid Gaussian random vector with zero mean, the averaged residuals $\bar{\epsilon}_i^\pm = \frac{1}{|\mathcal{A}^\pm|} \sum_i \epsilon_i^\pm$ are also Gaussian random vectors whose variances scale as $\frac{1}{|\mathcal{A}^\pm|}$, respectively. Since $|\mathcal{A}^+|$ and $|\mathcal{A}^-|$ are typically equal or at least of the same order, the residual difference $\Delta\epsilon = \bar{\epsilon}_i^+ - \bar{\epsilon}_i^-$ is also Gaussian with mean zero and variance scaling as $\frac{1}{|\mathcal{A}^+|}$. By standard concentration inequalities, this implies $\|\Delta\epsilon\|^2 = O(\frac{1}{|\mathcal{A}^+|})$ with high probability. Substituting this and the decomposition of the DiM direction $\mathbf{r}$ from (8) into $\cos(\mathbf{r}, \mathbf{r}_c) = \frac{\langle \mathbf{r}, \mathbf{r}_c \rangle}{\|\mathbf{r}\| \|\mathbf{r}_c\|}$, and using the assumption that the concept directions $\mathbf{r}_j$ are orthogonal with unit norm yields (9). $\square$

Lemma 1 assumes that the concept directions are orthogonal in order to simplify the analysis; in practice, many concepts are approximately orthogonal. When concept directions are not orthogonal, identifying the target concept direction becomes more complicated due to interactions between correlated concepts. Nevertheless, the main insight of the analysis still holds: ignoring the residual term, (9) shows that the alignment between the DiM direction $\mathbf{r}$ and the target concept direction $\mathbf{r}_c$ depends on two competing factors.

- **Target concept (signal):** the target concept difference $\Delta\alpha_c = \bar{\alpha}_c^+ - \bar{\alpha}_c^-$ should be as large as possible.
- **Non-target concepts (interference):** the differences associated with non-target concepts $\Delta\alpha_j = \bar{\alpha}_j^+ - \bar{\alpha}_j^-$, $j \neq c$ should be as small as possible.

From this perspective, averaging over many token activations primarily serves to cancel non-target concepts so that $\bar{\alpha}_j^+ \approx \bar{\alpha}_j^-$ for $j \neq c$, rather than directly increasing the target signal $\Delta\alpha_c$. However, transformer activations encode many simultaneous concepts and features, including topic, syntax, instruction format, stylistic patterns, and other context-dependent factors. Thus, a single prompt, or even a single token activation, may contain substantial non-target concept components. In practice,

the strengths of these non-target concepts are difficult to balance across $\mathcal{D}^+$ and $\mathcal{D}^-$. As a result, the cancellation condition $\bar{\alpha}_j^+ \approx \bar{\alpha}_j^-$ may fail even when a large number of examples are used. Moreover, increasing the number of construction examples can introduce additional stable non-target structure shared within each dataset. Consequently, averaging over more activations can reduce the random residual term $\Delta\epsilon$, while still preserving some systematic non-target concepts.

This perspective suggests that DiM construction should not be viewed purely as a variance-reduction problem through larger sample averages. Instead, it is more naturally viewed as a *robust contrastive mean estimation* problem, where the key challenge is to estimate the target concept direction while mitigating interference from stable but intervention-irrelevant concept directions. Before introducing new methods in the next subsection, we first use experiments to support this perspective.

Experimental Setup We evaluate DiM construction choices across two model families and three target concepts. The models are OLMo3-7B, OLMo3-32B [32], Qwen3-1.7B, and Qwen3-14B [33]. The target concepts are safety [11], politeness [34], and sentiment [9]. For each concept, construction and evaluation data are disjoint. Construction data are used only to estimate DiM directions, while held-out evaluation data are used to measure the effect of the resulting intervention. We evaluate safety on HarmBench, sentiment on SST-2, and politeness on PoliteGuard [35–37]. To isolate direction construction, we keep the downstream intervention protocol fixed across all construction rules. Thus, differences in target-concept score and general capability primarily reflect differences in the constructed direction rather than differences in how the direction is applied. Implementation details for the fixed intervention protocol are provided in Appendix A.

Since the concept directions $\mathbf{r}_j$ are unknown *a priori* and difficult to verify directly in practice, we report two complementary measurements. Directional stability is measured by the cosine similarity between DiM directions estimated from independent trials, indicating how reproducible a construction rule is under resampling. In addition, following prior work, we use downstream steering performance as a causal indicator of the quality of the identified target concept direction. We evaluate steering performance on the constructed target dataset as a measure of target steering effectiveness, and on the general benchmark MMLU-Pro [38] as a measure of the impact of steering on general capability.

Observation 1: Not All Prompt Tokens Are Equally Useful. Prior work often constructs steering directions from the final prompt token, while recent-token windows are also used as practical heuristics in representation steering and related intervention methods [39–43]. These choices implicitly assume that prompt-token position matters for DiM construction. We therefore ask whether averaging more prompt tokens yields a better direction, or whether different prompt positions contribute different amounts of intervention-relevant signal.

We compare DiM directions constructed from recent prompt-token windows and from all prompt tokens. For each token-selection rule, we fix the construction budget to 200 selected prompt-token activations per class. Thus, all rules use the same activation budget and differ only in prompt positions.

As shown in Figure 1a, token selection strongly affects both directional stability and downstream performance. All-token construction often gives weaker safety steering and larger general-capability degradation than recent-token construction. This suggests that earlier prompt tokens can introduce more task-irrelevant information, while later prompt tokens are closer to the generation boundary and more directly encode the information that conditions the target behavior. Thus, averaging more prompt tokens can add stable but intervention-irrelevant structure, rather than producing a more useful steering direction. Additional results for sentiment and politeness are provided in Appendix B.

Observation 2: Prompt Activations Are More Intervention-Relevant than Responses In addition to the input prompt, response tokens may contain continuation-specific information, such as stylistic patterns, generation dynamics, and other concepts expressed in the completion. Both prompt-only activations and prompt-response activations have been used for constructing DiM directions [9, 11, 12]. This raises a natural question: do response tokens provide a stronger concept signal for DiM construction, or do they instead introduce additional variation that is less useful for intervention? To answer this, we compare prompt-only and prompt-response construction under matched token-selection rules. Prompt-only directions use activations drawn only from the prompt, while prompt-response directions use activations drawn from the concatenated prompt and response. For each source, we consider the last token, the last five tokens, and all tokens, and we fix the construction budget to 200 selected activations per class for every source–selection combination.




Figure 1: **Diagnosing construction choices.** We vary three construction choices: token position, activation source, and example number. Panels report directional stability, safety steering performance, and general capability. For each setting, we run three independent construction trials and report the mean and standard deviation. Additional results for more concepts are provided in Appendix B.

Figure 1b shows that prompt-only directions generally achieve stronger safety steering under the same token-selection rule. Thus, tokens that explicitly express the target behavior are not necessarily the most useful tokens for constructing an intervention direction. A likely explanation is that response activations contain additional non-target structure from the generated continuation, which can dilute the prompt-level signal relevant for steering. Within each source, token choice remains important: all-token construction often underperforms more selective token choices. These results show that source choice is not a cosmetic implementation detail: DiM directions are strongest when constructed from activations that condition the target behavior, rather than from tokens that merely express it.

Observation 3: More Examples Improve Stability but not Steering We now turn to the question of whether increasing the number of contrastive examples will reduce sampling noise and yield a more reliable estimate of the target-concept direction. This question is also practically important because prior work uses widely varying construction-set sizes, from a few contrastive pairs to hundreds or thousands of examples [7–9, 12]. We vary the number of construction examples N while keeping the token-selection rule fixed to last-five prompt tokens, which yields a robust estimation according to previous experiments. For each N , we estimate directions from three independently sampled construction subsets and evaluate directional stability, steering performance, and general capability.

Figure 1c shows that increasing N generally improves directional stability, as directions become more similar across independent trials. However, this reproducibility does not consistently translate into better downstream steering. Safety steering is often strongest at small or medium construction sizes and can decrease as N grows, while general capability remains flat or declines. These results show that more examples reduce finite-sample variation but do not necessarily produce more useful intervention directions. Larger construction sets can make stable non-target structure more reproducible as well. Thus, the averaging problem cannot be solved by scale alone: effective DiM construction requires selecting intervention-relevant activations, not merely averaging more of them.




Figure 2: **Effect of Robust DiM construction.** We compare Robust DiM with a last-five-token DiM using the same candidate pool. Results are averaged over three independent runs. Robust DiM selects high-margin activations before computing the contrastive mean, yielding stronger safety steering and better general-capability preservation. Additional concept results are provided in Appendix B.

3.2 Robust DiM: Representative and Discriminative Activation Selection

The preceding analysis shows that DiM quality depends not only on how many activations are averaged, but also on which activations are used to estimate the class means. Importantly, the quality of the concept direction identified by DiM does not necessarily improve as the number of prompts increases. This supports our earlier view that identifying the target concept direction $\mathbf{r}_c$ from hidden states $\mathbf{h}_i^\pm$ should be treated as a robust contrastive mean estimation problem rather than a pure variance-reduction problem. Motivated by this perspective, we propose *Robust DiM*, which retains the simplicity of standard DiM while more reliably identifying intervention-relevant concept directions. Following the analysis in (9), the key idea behind Robust DiM is to preferentially select hidden states that exhibit a large target concept difference $\Delta\alpha_c$ while minimizing interference from non-target concepts, i.e., small $\Delta\alpha_j$ for $j \neq c$, by selecting those that are both representative of their own class and well separated from the opposite class.

Specifically, for each hidden state $\mathbf{h} \in \mathcal{A}^+$ (we illustrate the construction for the positive class; the negative class is handled analogously), we define the relative-margin score as

$$s^+(\mathbf{h}) = \frac{\|\mathbf{h} - \boldsymbol{\mu}^-\|_2^2 - \|\mathbf{h} - \boldsymbol{\mu}^+\|_2^2}{\|\mathbf{h} - \boldsymbol{\mu}^-\|_2^2 + \|\mathbf{h} - \boldsymbol{\mu}^+\|_2^2}. \quad (10)$$

This score is high when $\mathbf{h}$ is substantially closer to its own class center than to the opposite-class center. Consequently, high-scoring activations are both representative of their own class and contrastively separated from the opposing class. We then select the top- K high-margin activations under this score, denoted as $\tilde{\mathcal{A}}^+ = \text{TopK}_{\mathbf{h} \in \mathcal{A}^+} s^+(\mathbf{h})$. The resulting Robust DiM direction is computed as

$$\mathbf{r}_{\text{R-DiM}} = \frac{1}{|\tilde{\mathcal{A}}^+|} \sum_{\mathbf{h} \in \tilde{\mathcal{A}}^+} \mathbf{h} - \frac{1}{|\tilde{\mathcal{A}}^-|} \sum_{\mathbf{h} \in \tilde{\mathcal{A}}^-} \mathbf{h}, \quad \hat{\mathbf{r}}_{\text{R-DiM}} = \frac{\mathbf{r}_{\text{R-DiM}}}{\|\mathbf{r}_{\text{R-DiM}}\|_2}. \quad (11)$$

From the robust estimation perspective, this procedure acts as a filtering step that suppresses activations dominated by unstable or non-target concepts before mean estimation, analogous in spirit to filtering-based robust estimators such as trimmed means or median-of-means methods. Robust DiM differs from standard DiM only in which activations are used to estimate the class means. The intervention rule remains unchanged: $\hat{\mathbf{r}}_{\text{R-DiM}}$ can be used for additive steering to elicit the target behavior, or for directional ablation to suppress it.

Robust DiM identifies higher-quality concept directions that improve both target steering performance and general capability. We compare Robust DiM with a last-five-token DiM baseline. Both methods use the same candidate pool, consisting of 200 activations per class collected from the last five prompt tokens. The baseline averages all candidate activations, while Robust DiM selects the top $K = 50$ high-margin activations per class. All other settings are kept fixed.

Figure 2 shows that Robust DiM improves over the last-five-token DiM baseline across all four models. Under the same candidate pool, selecting high-margin activations yields stronger safety steering and better MMLU-Pro preservation than averaging all recent-token activations. This supports our central claim that steering vectors benefit from filtering construction activations rather than treating all selected prompt activations as equally useful. We also observe that larger models tend to preserve general capability better, but do not necessarily achieve higher safety steering scores, suggesting that scale alone does not remove the need for better construction-token selection.




Figure 3: **Robust DiM ablations.** On OLMo3-7B, we ablate the number of selected activations and the candidate-pool size. Results are averaged over three independent runs. Compared with using all candidate activations directly, selecting high-margin activations improves stability and downstream performance over averaging all candidates.




Figure 4: **Prefix-length ablation.** Using OLMo3-7B, we vary the number of generated tokens that receive the steering update. Results are averaged over three independent runs. **Left:** KL divergence from the unsteered model. **Right:** steering-capability trade-off. Short prefixes recover most of the all-token steering effect while inducing less distributional shift.

Ablation study. We next ablate the two selection choices in Robust DiM: the number of selected activations K and the number of candidate activations available before filtering. The first ablation varies K while keeping the candidate pool fixed. The second fixes $K = 50$ and varies the candidate-pool size. As shown in Figure 3(a), very small K gives unstable estimates, while very large K reintroduces lower-ranked activations and weakens the steering-capability trade-off. Compared with Figure 1(c), Figure 3(b) shows that Robust DiM achieves improved stability, steering performance, and general capability as the number of candidate activations increases. This contrasts with standard DiM, where steering performance degrades as more activations are averaged. These results suggest that filtering suppresses lower-quality activations and allows larger candidate pools to improve estimation without introducing as much non-target interference. In the following experiments, Robust DiM selects $K = 50$ from a candidate pool of 200 activations per class.

4 Prefix Steering: Early Tokens Are Sufficient for Control

The previous section studied which activations should be used for computing the steering direction. We now turn to the second token-level choice in activation steering: which generated token positions should receive the intervention. A common default is to apply the same steering update in (5) at every generated position k , implicitly treating all generated tokens as equally useful intervention sites [7, 11, 12]. However, while a steering update is intended to bias the model toward a target behavior, the same intervention can also interfere with unrelated computation, especially after the overall response trajectory has already been established. This raises an important question of whether steering should be applied uniformly across all generated tokens.




Figure 5: **All-token vs. prefix steering.** Results are averaged over three independent runs. Across models, five-token prefix steering retains most of the all-token intervention effect while better preserving MMLU-Pro performance. Additional results for more concepts are provided in Appendix B.

To answer this question, we apply the steering update only to an initial prefix of generated tokens (say prefix length m) and study how the number of steered tokens affects downstream behavior. Figure 4(a) shows the KL divergence between the steered and unsteered models across generation steps. On target concept tasks, the KL divergence increases sharply within the first few generated tokens and then grows much more slowly thereafter. In contrast, on the general benchmark MMLU-Pro, the KL divergence changes less during the early generation steps but continues to accumulate over later tokens. Consistent with this observation, Figure 4(b) shows that the target concept score rises rapidly within the first few generated tokens and quickly saturates near the all-token steering effect, while MMLU-Pro performance tends to degrade as the steering window becomes longer. These results suggest that early generated tokens capture most of the useful steering effect, whereas later interventions introduce additional distributional shift and interference with general capability.

This motivates our proposed **prefix steering**, which applies the steering update in (5) only to the first m generated tokens, leaving later tokens unmodified. All construction, layer, and steering-strength settings are kept fixed. Based on the above observations, we instantiate prefix steering using a fixed window of five generated tokens. This choice is not intended to optimize the prefix length for each individual model, but rather to test whether a simple, model-independent intervention policy can capture most of the benefits of full all-token steering. More adaptive strategies, such as dynamically selecting the steering window based on context or generation dynamics, may further improve performance and are left for future work.

Prefix steering preserves most steering benefits while better maintaining general capability.

Figure 5 shows the performance of prefix steering across different models. Prefix steering recovers most of the target-control effect achieved by full-generation intervention, while consistently preserving higher MMLU-Pro performance. This behavior is consistent with the KL analysis in Figure 4: short-prefix interventions remain much closer to the unsteered model, especially on general-task inputs, whereas longer intervention windows continue to accumulate distributional shift even after the target-control effect has largely saturated. Thus, prefix steering improves the steering—capability trade-off by concentrating the intervention on the early generated tokens that are most important for shaping the response trajectory, while avoiding unnecessary perturbations at later generation steps.

5 Conclusion

We studied DiM steering as a token-level robust contrastive mean estimation problem. Both theoretical analysis and empirical results demonstrate that increasing the number of activations, either by using more prompt and response tokens or by increasing the number of contrastive examples, does not necessarily improve steering performance and can even degrade it. We identify the key challenge as interference from stable but intervention-irrelevant concepts encoded in hidden states, motivating a reinterpretation of DiM as a robust contrastive mean estimation problem.

We introduced Robust DiM, which constructs directions from high-margin activations that are both class-representative and contrastively separated. Across three target concepts and four models, Robust DiM improves target-concept steering while improving general-capability preservation. We further introduced prefix steering, which applies the intervention only to early generated tokens, which recovers most of the target-control effect of full-generation intervention while substantially reducing interference with general capabilities.

References

- [1] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, Sandhini Agarwal, Ariel Herbert-Voss, Gretchen Krueger, Tom Henighan, Rewon Child, Aditya Ramesh, Daniel Ziegler, Jeffrey Wu, Clemens Winter, Chris Hesse, Mark Chen, Eric Sigler, Mateusz Litwin, Scott Gray, Benjamin Chess, Jack Clark, Christopher Berner, Sam McCandlish, Alec Radford, Ilya Sutskever, and Dario Amodei. Language models are few-shot learners. In H. Larochelle, M. Ranzato, R. Hadsell, M.F. Balcan, and H. Lin, editors, *Advances in Neural Information Processing Systems*, volume 33, pages 1877–1901. Curran Associates, Inc., 2020. URL https://proceedings.neurips.cc/paper_files/paper/2020/file/1457c0d6bfc4967418bfb8ac142f64a-Paper.pdf.
- [2] Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, Charles Sutton, Sebastian Gehrmann, et al. Palm: Scaling language modeling with pathways. *Journal of machine learning research*, 24(240): 1–113, 2023.
- [3] Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- [4] Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual associations in gpt. *Advances in neural information processing systems*, 35:17359–17372, 2022.
- [5] Nelson Elhage, Neel Nanda, Catherine Olsson, Tom Henighan, Nicholas Joseph, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, Nova DasSarma, Dawn Drain, Deep Ganguli, Zac Hatfield-Dodds, Danny Hernandez, Andy Jones, Jackson Kernion, Liane Lovitt, Kamal Ndousse, Dario Amodei, Tom Brown, Jack Clark, Jared Kaplan, Sam McCandlish, and Chris Olah. A mathematical framework for transformer circuits. *Transformer Circuits Thread*, 2021. <https://transformer-circuits.pub/2021/framework/index.html>.
- [6] Rishi Bommasani, Drew A Hudson, Ehsan Adeli, Russ Altman, Simran Arora, Sydney von Arx, Michael S Bernstein, Jeannette Bohg, Antoine Bosselut, Emma Brunskill, et al. On the opportunities and risks of foundation models. *arXiv preprint arXiv:2108.07258*, 2021.
- [7] Alexander Matt Turner, Lisa Thiergart, Gavin Leech, David Udell, Juan J Vazquez, Ulisse Mini, and Monte MacDiarmid. Steering language models with activation engineering. *arXiv preprint arXiv:2308.10248*, 2023.
- [8] Kenneth Li, Oam Patel, Fernanda Viégas, Hanspeter Pfister, and Martin Wattenberg. Inference-time intervention: Eliciting truthful answers from a language model. *Advances in Neural Information Processing Systems*, 36:41451–41530, 2023.
- [9] Andy Zou, Long Phan, Sarah Chen, James Campbell, Phillip Guo, Richard Ren, Alexander Pan, Xuwang Yin, Mantas Mazeika, Ann-Kathrin Dombrowski, et al. Representation engineering: A top-down approach to ai transparency. *arXiv preprint arXiv:2310.01405*, 2023.
- [10] Alex Mallen, Madeline Brumley, Julia Kharchenko, and Nora Belrose. Eliciting latent knowledge from quirky language models. *arXiv preprint arXiv:2312.01037*, 2023.
- [11] Andy Ardit, Oscar Obeso, Aaqib Syed, Daniel Paleka, Nina Panickssery, Wes Gurnee, and Neel Nanda. Refusal in language models is mediated by a single direction. *Advances in Neural Information Processing Systems*, 37:136037–136083, 2024.
- [12] Nina Rimskey, Nick Gabrieli, Julian Schulz, Meg Tong, Evan Hubinger, and Alexander Turner. Steering llama 2 via contrastive activation addition. In Lun-Wei Ku, Andre Martins, and Vivek Srikumar, editors, *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15504–15522, Bangkok, Thailand, August 2024. Association for Computational Linguistics. doi: 10.18653/v1/2024.acl-long.828. URL <https://aclanthology.org/2024.acl-long.828/>.

- [13] Zhengxuan Wu, Aryaman Arora, Atticus Geiger, Zheng Wang, Jing Huang, Dan Jurafsky, Christopher D Manning, and Christopher Potts. Axbench: Steering llms? even simple baselines outperform sparse autoencoders. *arXiv preprint arXiv:2501.17148*, 2025.
- [14] Daniel Tan, David Chanin, Aengus Lynch, Brooks Paige, Dimitrios Kanoulas, Adrià Garriga-Alonso, and Robert Kirk. Analysing the generalisation and reliability of steering vectors. *Advances in Neural Information Processing Systems*, 37:139179–139212, 2024.
- [15] Yixiong Hao, Ayush Panda, Stepan Shabalin, and Sheikh Abdur Raheem Ali. Patterns and mechanisms of contrastive activation engineering. *arXiv preprint arXiv:2505.03189*, 2025.
- [16] Joschka Braun, Carsten Eickhoff, David Krueger, Seyed Ali Bahrainian, and Dmitrii Krasheninnikov. Understanding (un) reliability of steering vectors in language models. *arXiv preprint arXiv:2505.22637*, 2025.
- [17] Navita Goyal and Hal Daumé III. Steering safely or off a cliff? rethinking specificity and robustness in inference-time interventions. *arXiv preprint arXiv:2602.06256*, 2026.
- [18] Kiho Park, Yo Joong Choe, and Victor Veitch. The linear representation hypothesis and the geometry of large language models. *arXiv preprint arXiv:2311.03658*, 2023.
- [19] Wes Gurnee and Max Tegmark. Language models represent space and time. *arXiv preprint arXiv:2310.02207*, 2023.
- [20] Curt Tigges, Oskar John Hollinsworth, Atticus Geiger, and Neel Nanda. Linear representations of sentiment in large language models. *arXiv preprint arXiv:2310.15154*, 2023.
- [21] Peter J Liu, Mohammad Saleh, Etienne Pot, Ben Goodrich, Ryan Sepassi, Lukasz Kaiser, and Noam Shazeer. Generating wikipedia by summarizing long sequences. *arXiv preprint arXiv:1801.10198*, 2018.
- [22] Joris Postmus and Steven Abreu. Steering large language models using conceptors: Improving addition-based activation engineering. *arXiv preprint arXiv:2410.16314*, 2024.
- [23] Shivam Raval, Hae Jin Song, Linlin Wu, Abir Harrasse, Jeff Phillips, and Amirali Abdullah. Curveball steering: The right direction to steer isn’t always linear. *arXiv preprint arXiv:2603.09313*, 2026.
- [24] Nora Belrose, David Schneider-Joseph, Shauli Ravfogel, Ryan Cotterell, Edward Raff, and Stella Biderman. Leace: Perfect linear concept erasure in closed form. *Advances in Neural Information Processing Systems*, 36:66044–66063, 2023.
- [25] Hoagy Cunningham, Aidan Ewart, Logan Riggs, Robert Huben, and Lee Sharkey. Sparse autoencoders find highly interpretable features in language models. *arXiv preprint arXiv:2309.08600*, 2023.
- [26] Been Kim, Martin Wattenberg, Justin Gilmer, Carrie Cai, James Wexler, Fernanda Viegas, et al. Interpretability beyond feature attribution: Quantitative testing with concept activation vectors (tcav). In *International conference on machine learning*, pages 2668–2677. PMLR, 2018.
- [27] Tsung-Min Pai, Jui-I Wang, Li-Chun Lu, Shao-Hua Sun, Hung-Yi Lee, and Kai-Wei Chang. Billy: Steering large language models via merging persona vectors for creative generation. In *Conference of the European Chapter of the Association for Computational Linguistics*, 2026.
- [28] Weixuan Wang, JINGYUAN YANG, and Wei Peng. Semantics-adaptive activation intervention for LLMs via dynamic steering vectors. In *The Thirteenth International Conference on Learning Representations*, 2025. URL <https://openreview.net/forum?id=8WQ7VTfPTL>.
- [29] Samuel Soo, Chen Guang, Wesley Teng, Chandrasekaran Balaganesh, Tan Guoxian, and Yan Ming. Interpretable steering of large language models with feature guided activation additions. *arXiv preprint arXiv:2501.09929*, 2025.

- [30] Pyae Phoo Min, Avigya Paudel, Naufal Adityo, Arthur Zhu, Andrew Rufail, Cole Blondin, Kevin Zhu, Sunishchal Dev, and Sean O’Brien. Mitigating sycophancy in language models via sparse activation fusion and multi-layer activation steering. In *Mechanistic Interpretability Workshop at NeurIPS 2025*, 2025. URL <https://openreview.net/forum?id=BCS7HHInC2>.
- [31] Samuel Marks and Max Tegmark. The geometry of truth: Emergent linear structure in large language model representations of true/false datasets. *arXiv preprint arXiv:2310.06824*, 2023.
- [32] Team Olmo, Allyson Ettinger, Amanda Bertsch, Bailey Kuehl, David Graham, David Heineman, Dirk Groeneveld, Faeze Brahman, Finbarr Timbers, Hamish Ivison, et al. Olmo 3. *arXiv preprint arXiv:2512.13961*, 2025.
- [33] An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.
- [34] Cristian Danescu-Niculescu-Mizil, Moritz Sudhof, Dan Jurafsky, Jure Leskovec, and Christopher Potts. A computational approach to politeness with application to social factors. In Hinrich Schuetze, Pascale Fung, and Massimo Poesio, editors, *Proceedings of the 51st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 250–259, Sofia, Bulgaria, August 2013. Association for Computational Linguistics. URL <https://aclanthology.org/P13-1025/>.
- [35] Mantas Mazeika, Long Phan, Xuwang Yin, Andy Zou, Zifan Wang, Norman Mu, Elham Sakhaee, Nathaniel Li, Steven Basart, Bo Li, et al. Harmbench: A standardized evaluation framework for automated red teaming and robust refusal. *arXiv preprint arXiv:2402.04249*, 2024.
- [36] Richard Socher, Alex Perelygin, Jean Wu, Jason Chuang, Christopher D. Manning, Andrew Ng, and Christopher Potts. Recursive deep models for semantic compositionality over a sentiment treebank. In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing*, pages 1631–1642, Seattle, Washington, USA, October 2013. Association for Computational Linguistics. URL <https://www.aclweb.org/anthology/D13-1170>.
- [37] Intel. Politeguard. Hugging Face dataset, 2025. Available at <https://huggingface.co/datasets/Intel/polite-guard>.
- [38] Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, Abhranil Chandra, Shiguang Guo, Weiming Ren, Aaran Arulraj, Xuan He, Ziyang Jiang, et al. Mmlu-pro: A more robust and challenging multi-task language understanding benchmark. *Advances in Neural Information Processing Systems*, 37:95266–95290, 2024.
- [39] Felix Hofstätter, Teun van der Weij, Jayden Teoh, Rada Djoneva, Henning Bartsch, and Francis Rhys Ward. The elicitation game: Evaluating capability elicitation techniques. In *Forty-second International Conference on Machine Learning*, 2025. URL <https://openreview.net/forum?id=kT0EVqL77E>.
- [40] Xinyan Jiang, Lin Zhang, Jiayi Zhang, Qingsong Yang, Guimin Hu, Di Wang, and Lijie Hu. Msrs: Adaptive multi-subspace representation steering for attribute alignment in large language models. *arXiv preprint arXiv:2508.10599*, 2025.
- [41] Minjun Zhu, Yixuan Weng, Linyi Yang, Yifan Wei, Ningyu Zhang, and Yue Zhang. Locking down the finetuned LLMs safety, 2025. URL <https://openreview.net/forum?id=YGoF15KKFc>.
- [42] Hannah Cyberek, Yangfeng Ji, and David Evans. Unsupervised concept vector extraction for bias control in llms. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 28321–28343, 2025.
- [43] Chen Qian, Dongrui Liu, Jie Zhang, Yong Liu, and Jing Shao. DEAN: Deactivating the coupled neurons to mitigate fairness-privacy conflicts in large language models, 2024. URL <https://openreview.net/forum?id=v4PnwdA056>.

- [44] Abhay Sheshadri, Aidan Ewart, Phillip Guo, Aengus Lynch, Cindy Wu, Vivek Hebbar, Henry Sleight, Asa Cooper Stickland, Ethan Perez, Dylan Hadfield-Menell, et al. Latent adversarial training improves robustness to persistent harmful behaviors in llms. *arXiv preprint arXiv:2407.15549*, 2024.
- [45] Lianmin Zheng, Wei-Lin Chiang, Ying Sheng, Siyuan Zhuang, Zhanghao Wu, Yonghao Zhuang, Zi Lin, Zhuohan Li, Dacheng Li, Eric Xing, et al. Judging llm-as-a-judge with mt-bench and chatbot arena. *Advances in neural information processing systems*, 36:46595–46623, 2023.
- [46] Nishant Subramani, Nivedita Suresh, and Matthew E Peters. Extracting latent steering vectors from pretrained language models. In *Findings of the Association for Computational Linguistics: ACL 2022*, pages 566–581, 2022.
- [47] Sumanth Dathathri, Andrea Madotto, Janice Lan, Jane Hung, Eric Frank, Piero Molino, Jason Yosinski, and Rosanne Liu. Plug and play language models: A simple approach to controlled text generation. *arXiv preprint arXiv:1912.02164*, 2019.
- [48] Hanyu Zhang, Xiting Wang, Chengao Li, Xiang Ao, and Qing He. Controlling large language models through concept activation vectors. In *Proceedings of the AAAI Conference on Artificial Intelligence*, volume 39, pages 25851–25859, 2025.
- [49] Alessandro Stolfo, Vidhisha Balachandran, Safoora Yousefi, Eric Horvitz, and Besmira Nushi. Improving instruction-following in language models through activation steering. *arXiv preprint arXiv:2410.12877*, 2024.
- [50] Dana Arad, Aaron Mueller, and Yonatan Belinkov. Saes are good for steering—if you select the right features. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, pages 10252–10270, 2025.
- [51] Hannah Sterz, Fabian David Schmidt, Goran Glavaš, and Ivan Vulic. Recover the target language: Language steering without sacrificing task performance. *Preprint*, 2025.
- [52] Jinghao Zhang, Yuting Liu, Wenjie Wang, Qiang Liu, Shu Wu, Liang Wang, and Tat-Seng Chua. Personalized text generation with contrastive activation steering. In Wanxiang Che, Joyce Nabende, Ekaterina Shutova, and Mohammad Taher Pilehvar, editors, *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 7128–7141, Vienna, Austria, July 2025. Association for Computational Linguistics. ISBN 979-8-89176-251-0. doi: 10.18653/v1/2025.acl-long.353. URL <https://aclanthology.org/2025.acl-long.353/>.
- [53] Kai Konen, Sophie Jentzsch, Diaoulé Diallo, Peer Schütt, Oliver Bensch, Roxanne El Baff, Dominik Opitz, and Tobias Hecking. Style vectors for steering generative large language models. In *Findings of the Association for Computational Linguistics: EACL 2024*, pages 782–802, 2024.
- [54] Xudong Zhu, Jiachen Jiang, Mohammad Mahdi Khalili, and Zhihui Zhu. From emergence to control: Probing and modulating self-reflection in language models, 2026. URL <https://openreview.net/forum?id=79uLTCNmpJ>.
- [55] Yi Fang, Wenjie Wang, Mingfeng Xue, Boyi Deng, Fengli Xu, Dayiheng Liu, and Fuli Feng. Controllable llm reasoning via sparse autoencoder-based steering. *arXiv preprint arXiv:2601.03595*, 2026.
- [56] Peixuan Han, Cheng Qian, Xiusi Chen, Yuji Zhang, Denghui Zhang, and Heng Ji. Safeswitch: Steering unsafe llm behavior via internal activation signals. *arXiv preprint arXiv:2502.01042*, 2025.
- [57] Xudong Zhu, Mohammad Mahdi Khalili, and Zhihui Zhu. Abstopk: Rethinking sparse autoencoders for bidirectional features. In *The Fourteenth International Conference on Learning Representations*, 2026. URL <https://openreview.net/forum?id=EEs6I4c07S>.

- 535 [58] Zara Siddique, Irtaza Khalid, Liam Turner, and Luis Espinosa-Anke. Shifting perspectives:
536 Steering vectors for robust bias mitigation in LLMs. In Vera Demberg, Kentaro Inui, and Lluís
537 Marquez, editors, *Findings of the Association for Computational Linguistics: EACL 2026*, pages
538 809–820, Rabat, Morocco, March 2026. Association for Computational Linguistics. ISBN
539 979-8-89176-386-9. doi: 10.18653/v1/2026.findings-eacl.41. URL <https://aclanthology.org/2026.findings-eacl.41/>.
- 541 [59] Xuanqi Zhang, Haoyang Shang, and Xiaoxiao Li. Gss: Gated subspace steering for selective
542 memorization mitigation in llms. *arXiv preprint arXiv:2602.08901*, 2026.
- 543 [60] Van Dai Do, Quan Hung Tran, Svetha Venkatesh, and Hung Le. Dynamic steering with
544 episodic memory for large language models. In Wanxiang Che, Joyce Nabende, Ekaterina
545 Shutova, and Mohammad Taher Pilehvar, editors, *Findings of the Association for Computational
546 Linguistics: ACL 2025*, pages 13731–13749, Vienna, Austria, July 2025. Association for
547 Computational Linguistics. ISBN 979-8-89176-256-5. doi: 10.18653/v1/2025.findings-acl.706.
548 URL <https://aclanthology.org/2025.findings-acl.706/>.

A Experimental Details

Dataset. For each concept, we use separate construction and evaluation data.

- **Safety.** We construct the safety direction using LLM-LAT/benign-dataset as the positive set and LLM-LAT/harmful-dataset as the negative set [44], and evaluate downstream safety behavior on HarmBench [35].
- **Sentiment.** We construct sentiment directions using the SST-2 training split and evaluate on the SST-2 validation split [36].
- **Politeness.** We construct politeness directions using the PoliteGuard training split and evaluate on the PoliteGuard test split [37]. We restrict the dataset to the two endpoint labels, *polite* and *impolite*, and exclude intermediate politeness labels.

All reported concept scores are computed on evaluation examples that are disjoint from the examples used to construct the DiM directions.

Generation and evaluation protocol. For target-concept evaluation, we apply interventions during autoregressive generation and use greedy decoding with a maximum generation length of 512 tokens, following HarmBench [35]. Decoding settings are fixed across different settings.

For safety, we evaluate elicitation on HarmBench prompts using the HarmBench evaluation protocol. Since HarmBench consists of harmful-behavior prompts rather than a paired positive-negative classification dataset, we use it only to measure whether additive steering increases safe behavior on these prompts; we do not evaluate a corresponding suppression setting for safety. For sentiment and politeness, we evaluate both elicitation and suppression on held-out labeled examples. Elicitation is evaluated by applying additive steering to negative-class inputs. Suppression is evaluated by applying directional ablation to positive-class inputs.

To measure capability preservation, we evaluate the intervened model on MMLU-Pro [38] using the standard multiple-choice evaluation protocol. We refer to the resulting score as the general score, which measures whether the intervention preserves general task performance.

When constructing prompt-response directions, we collect activations from the concatenated prompt and model response. For all-token prompt-response construction, we cap the response length at 50 generated tokens before collecting activations. This prevents a small number of long responses from dominating the activation budget and ensures that each direction is constructed from activations drawn from multiple examples.

Intervention layer. For each model-concept pair, we select the intervention layer on a held-out tuning split and then keep it fixed for all reported evaluation experiments. The same selected layer is used for all construction rules and intervention policies within that model-concept pair. This avoids method-specific layer tuning and ensures that comparisons isolate the effect of direction construction and intervention placement.

Intervention strength. We scale the additive intervention coefficient by the typical norm of the model’s hidden activations. For each model and intervention layer, let $\bar{r}$ denote the average ℓ_2 norm of the residual-stream activations at that layer. We evaluate $\alpha \in \{0.01\bar{r}, 0.1\bar{r}, \bar{r}\}$. For scalar summaries, we select the coefficient that achieves the highest concept score among coefficients whose general score remains within 90% of the unsteered model’s general score. This criterion favors strong target-behavior control while excluding interventions that substantially degrade general capability. All methods being compared under a given setting use the same rule. For directional ablation experiments, we use $\beta \in \{0.01, 0.1, 1\}$ and apply the same coefficient rule to experiments.

A.1 Concept-Score Evaluation

For safety evaluation, we follow the standard HarmBench evaluation protocol [35], which allows direct comparison with prior activation-steering methods. For sentiment and politeness, we use an LLM-as-a-judge protocol to classify each generated response [45], with GPT-5.5 as the judge model. The judge receives the input prompt and the model response and uses the same judge prompt for all methods within each concept.

For both sentiment and politeness, the judge assigns one of three labels: the positive endpoint, the negative endpoint, or NEUTRAL_OR_MIXED. The positive endpoint is POSITIVE for sentiment and POLITE for politeness. For elicitation, the concept score is the fraction of responses labeled with the positive endpoint, while for suppression, the concept score is the fraction of responses not labeled with the positive endpoint. Therefore, suppression is evaluated as removal of the target behavior, rather than forced generation of the opposite behavior.

Sentiment judge prompt.

Task. Classify the sentiment of the assistant response as one of POSITIVE, NEGATIVE, or NEUTRAL_OR_MIXED.

Rules. Judge the sentiment expressed by the assistant response itself, not the sentiment of the user prompt. Use POSITIVE for responses expressing approval, praise, satisfaction, optimism, or favorable evaluation. Use NEGATIVE for responses expressing criticism, dissatisfaction, pessimism, dislike, or unfavorable evaluation. Use NEUTRAL_OR_MIXED when the response is neutral, ambiguous, mixed, off-topic, or does not express a clear sentiment.

Return only valid JSON:

```
{"label": "POSITIVE" | "NEGATIVE" | "NEUTRAL_OR_MIXED"}
```

User prompt: {prompt}

Assistant response: {response}

Politeness judge prompt.

Task. Classify the politeness of the assistant response as one of: POLITE, IMPOLITE, or NEUTRAL_OR_MIXED.

Instructions. Judge only the tone of the assistant response.

- Use POLITE for respectful, courteous, or tactful responses.
- Use IMPOLITE for rude, hostile, dismissive, insulting, or unnecessarily blunt responses.
- Use NEUTRAL_OR_MIXED when the tone is neutral, ambiguous, mixed, off-topic, or too short to determine.

Output format. Return only valid JSON:

```
{"label": "POLITE" | "IMPOLITE" | "NEUTRAL_OR_MIXED"}
```

User prompt: {prompt}

Assistant response: {response}

B Additional Evidence Across Target Concepts

To test whether the trends observed for safety are concept-specific, we repeat the same construction and intervention analyses for sentiment and politeness. For each concept, we follow the same evaluation protocol as in the main experiments: construction data are kept disjoint from evaluation data, all comparisons use matched activation budgets where applicable, and each setting is run with three independently sampled construction sets.

Figures 6 and 7 show that the main conclusions extend beyond safety. Across both sentiment and politeness, DiM performance remains sensitive to which activations enter the contrastive mean. Recent prompt-token directions often outperform all-token directions under the same activation budget, and prompt-only construction generally matches or improves over prompt-response construction. These patterns support the view that additional token activations are not automatically beneficial: they can also introduce stable, concept-irrelevant structure into the estimated direction. Similarly, increasing the number of construction examples improves directional reproducibility, but does not consistently improve steering performance or general-capability preservation.

The two proposed interventions show the same qualitative behavior on these additional concepts. Robust DiM improves over the last-five-token DiM baseline by selecting more representative and discriminative construction activations. Prefix steering, using only the first five generated tokens, recovers most of the target-control effect of all-token intervention while better preserving MMLU-Pro

626 performance. Together, these results indicate that both token selection during direction construction
627 and token placement during intervention are general design choices for DiM steering, rather than
628 artifacts of the safety setting.

629 C Related Works

630 **Activation steering and Difference-in-Means directions.** A growing line of work studies how
631 model behavior can be controlled by intervening directly on internal activations rather than modifying
632 model weights [46–50]. These methods use directions in activation space to elicit, suppress, or modify
633 behaviors at inference time [51–53]. Among them, DiM is a simple and widely used approach: given
634 positive and negative examples of a target behavior, it subtracts the mean activation of the negative
635 set from the mean activation of the positive set and uses the resulting direction for steering [27–29].
636 DiM-style directions have been applied to behaviors such as sentiment, refusal, truthfulness, safety,
637 and reasoning-related patterns [54–56]. Their appeal comes from requiring only contrastive examples,
638 without training an additional model or solving an optimization problem. However, most prior work
639 focuses on whether such directions can control behavior, while the practical choices that determine
640 how the direction is estimated and applied are often left implicit.

641 **Token granularity in direction construction and intervention.** Although DiM is usually pre-
642 sented as an averaging procedure, the average must be taken over a particular set of token activations.
643 Existing steering methods instantiate this choice in different ways, including using final-token activa-
644 tions, recent-token activations, all-token averages, prompt-only activations, or activations collected
645 from both prompts and responses [15, 57, 58]. A related choice appears during intervention: many
646 methods apply the same steering direction at every generated position, while adaptive or dynamic
647 steering methods suggest that the necessary intervention strength may vary over the course of genera-
648 tion [28, 59, 60]. These token-level choices affect the granularity at which steering operates, but they
649 are usually treated as implementation details rather than as objects of study. Our work makes this
650 granularity explicit by treating DiM as a token-level robust contrastive mean estimation problem, in
651 which token selection affects both how the direction is estimated and how the intervention is applied.

652 D Limitations

653 Our study focuses on DiM-style linear activation steering. Although Robust DiM improves con-
654 trastive direction construction, a single linear direction may not fully capture behaviors that are
655 multidimensional, nonlinear, or strongly context-dependent. Extending token-selective construction
656 to multi-directional steering methods remains an important direction for future work.

657 Our experiments cover four models from two model families and three target concepts: safety,
658 sentiment, and politeness. These settings provide evidence that the observed token-level effects are
659 not specific to a single model or behavior. However, they do not cover long-horizon agentic behavior,
660 multilingual generation, multi-turn dialogue, or settings where the target behavior changes over time.
661 The effectiveness of Robust DiM and prefix steering may therefore depend on the structure of the
662 target concept and the quality of the contrastive construction data.

663 Our general-capability evaluation uses MMLU-Pro as the main preservation metric. While MMLU-
664 Pro provides a useful measure of broad multiple-choice performance, it does not capture all possible
665 forms of capability degradation, such as changes in factuality, calibration, reasoning style, instruction
666 following, or long-form generation quality. In addition, our sentiment and politeness evaluations rely
667 on an LLM-as-a-judge protocol, which may introduce judge-specific biases despite using the same
668 judge prompt across methods.

669 E Broader impacts


670 This work studies inference-time activation steering, aiming to improve behavioral control while
671 better preserving general capabilities. More precise direction construction and localized interventions
672 may make steering methods easier to evaluate and less disruptive, especially for safety-relevant
673 behaviors that need to be encouraged or suppressed without retraining.

674 Activation steering is a general model-control technique, so safeguards are important when applying
675 it beyond controlled experiments. Steering directions should be built from explicit and auditable
676 contrastive datasets, and interventions should be evaluated for both target-behavior control and
677 unintended side effects. Localized policies such as prefix steering may help reduce broad distributional
678 shifts by limiting interventions to the generated positions where they are most useful. In deployment
679 settings, these methods should be paired with task-specific safety evaluations, monitoring, and
680 appropriate access controls.


681 **F LLM Usage**

682 We used LLMs to assist with language editing, writing clarity, and formatting during manuscript
683 preparation. We also used LLMs to help draft, refactor, and debug portions of the experimental and
684 plotting code. All generated text and code were reviewed and edited by the authors.


685 In addition, our sentiment and politeness evaluations use an LLM-as-a-judge protocol, with GPT-
686 5.5 as the judge model. The judge is used only to assign predefined concept labels to generated
687 responses under fixed prompts shared across all methods within each concept. This usage is part of
688 the evaluation protocol and is described in Section A.1.




(a) Varying prompt-token selection rules under a matched activation budget.




(b) Comparing prompt-only and prompt-response construction.



(c) Varying the number of construction examples N .




(d) Comparing Robust DiM with last-five-token DiM.




(e) Comparing all-token intervention with prefix intervention.


Figure 6: Additional sentiment results. We report sentiment counterparts of the main construction and intervention analyses. The panels follow the same order as in the main text: prompt-token selection, activation source, number of construction examples, Robust DiM construction, and intervention placement. These results further support that effective DiM steering depends on which token activations define the direction and which generated token positions receive the intervention.




(a) Varying prompt-token selection rules under a matched activation budget.




(b) Comparing prompt-only and prompt-response construction.



(c) Varying the number of construction examples N .



(d) Comparing Robust DiM with last-five-token DiM.



(e) Comparing all-token intervention $\mathcal{I}_{\text{all}}$ with prefix intervention $\mathcal{I}_5$.

Figure 7: Additional politeness results. We repeat the main construction and intervention analyses for politeness, following the same order as in the main text. The results show the same qualitative pattern: DiM depends strongly on which token activations define the direction and which generated token positions receive the intervention.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The claims are evaluateded across multiple models, concepts, and held-out datasets.

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes].

Justification: The paper includes a separate Limitations section in Appendix D.

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [N/A].

Justification: The paper does not present formal theoretical results or proofs. The mathematical discussion of DiM and the averaging assumption is used to motivate the empirical study rather than to establish a theorem.

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [Yes].

Justification: The paper specifies evaluation metrics in the main text and Appendix A.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [Yes]

Justification: We release anonymized code and scripts for reproducing the main experiments.

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).
- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: [Yes].

Justification: The paper describes the testing details. Since the method does not train model weights, training details are not applicable.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: [Yes].

Justification: For stability experiments, the paper reports the mean and standard deviation of cosine similarity across three trials. For downstream steering experiments, the paper reports results across multiple models and concepts; when error bars are shown, they correspond to variability across repeated construction trials.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The authors should answer [Yes] if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.

- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: [Yes].

Justification: The paper describes the experimental protocol in main contents and appendix.

Guidelines:

- The answer [N/A] means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.
- The paper should disclose whether the full research project required more compute than the experiments reported in the paper (e.g., preliminary or failed experiments that didn't make it into the paper).

9. Code of ethics

Question: Does the research conducted in the paper conform, in every respect, with the NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?

Answer: [Yes].

Justification: The work studies representation-level steering methods using publicly available models and datasets, and does not involve private data.

Guidelines:

- The answer [N/A] means that the authors have not reviewed the NeurIPS Code of Ethics.
- If the authors answer [No], they should explain the special circumstances that require a deviation from the Code of Ethics.
- The authors should make sure to preserve anonymity (e.g., if there is a special consideration due to laws or regulations in their jurisdiction).

10. Broader impacts

Question: Does the paper discuss both potential positive societal impacts and negative societal impacts of the work performed?

Answer: [Yes].

Justification: The paper discusses broader impacts in Appendix E.

Guidelines:

- The answer [N/A] means that there is no societal impact of the work performed.

- If the authors answer [N/A] or [No], they should explain why their work has no societal impact or why the paper does not address societal impact.
- Examples of negative societal impacts include potential malicious or unintended uses (e.g., disinformation, generating fake profiles, surveillance), fairness considerations (e.g., deployment of technologies that could make decisions that unfairly impact specific groups), privacy considerations, and security considerations.
- The conference expects that many papers will be foundational research and not tied to particular applications, let alone deployments. However, if there is a direct path to any negative applications, the authors should point it out. For example, it is legitimate to point out that an improvement in the quality of generative models could be used to generate Deepfakes for disinformation. On the other hand, it is not needed to point out that a generic algorithm for optimizing neural networks could enable people to train models that generate Deepfakes faster.
- The authors should consider possible harms that could arise when the technology is being used as intended and functioning correctly, harms that could arise when the technology is being used as intended but gives incorrect results, and harms following from (intentional or unintentional) misuse of the technology.
- If there are negative societal impacts, the authors could also discuss possible mitigation strategies (e.g., gated release of models, providing defenses in addition to attacks, mechanisms for monitoring misuse, mechanisms to monitor how a system learns from feedback over time, improving the efficiency and accessibility of ML).

11. Safeguards

Question: Does the paper describe safeguards that have been put in place for responsible release of data or models that have a high risk for misuse (e.g., pre-trained language models, image generators, or scraped datasets)?

Answer: [N/A].

Justification: The paper does not release a new pretrained model, image generator, scraped dataset, or other high-risk asset. The experiments use existing publicly available models and datasets.

Guidelines:

- The answer [N/A] means that the paper poses no such risks.
- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: [Yes].

Justification: We use these assets according to their intended research usage and will include license information where available.

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.

- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: [N/A].

Justification: The paper does not introduce a new dataset, benchmark, pretrained model, or other standalone asset.

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: [N/A].

Justification: The paper does not involve crowdsourcing experiments or human-subject data collection.

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [N/A].

Justification: The paper does not involve research with human subjects, crowdsourced participants, or newly collected human data. Therefore, IRB approval is not applicable.

Guidelines:

- 999
- 1000
- 1001
- 1002
- 1003
- 1004
- 1005
- 1006
- 1007
- 1008
- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
 - Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
 - We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
 - For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

1009

16. Declaration of LLM usage

1010

1011

1012

1013

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

1014

Answer: [Yes].

1015

Justification: The paper describes LLM usage in Appendix F and Section A.1.

1016

Guidelines:

- 1017
- 1018
- 1019
- 1020
- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
 - Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.